

CS4533 Concurrent Programming

Take Home Lab 1

Group Members:

- 200081B - Chandeepea E.A.C.
- 200698X - Weerasekara W.M.T.B.

Step 1

Design Explanation

Objective:

The first step in this lab involves designing a solution to generate operations for the linked list based on specified fractions (mMember, mInsert, mDelete) using a given number of threads. This design must efficiently distribute the operations among threads and accurately measure the execution time of these operations.

Files Structure:

linked_list.c and linked_list.h : Implementation of the linked list data structure.
serial.c : Source code for the serial version.
mutex.c : Source code using mutex locks.
rw_lock.c : Source code using read-write locks.
Makefile : Automates the build and run processes.

Explanation:

First, the linked list is initialized by n random values, and then the m number of operations is distributed among the number of threads equally. For instance, if there are 10,000 operations (m=10000) and 8 threads (thread_count=8), each thread will handle 1,250 operations so that each thread works on a specific set of operations.

Inside the ThreadFunction, it runs Member operations (m/thread_count)*mMember times, Insert operations (m/thread_count)*mInsert times, and Delete operations (m/thread_count)*mDelete times. ThreadFunction runs Member, Insert, and Delete operations in a random order. Used Pthreads library for mutex and read-write lock implementations.

- For the serial implementation, no synchronization is required.
- For the mutex-based parallel implementation, a single mutex lock controls access to the linked list, ensuring that only one thread can modify the list at any time.
- For the read-write lock implementation, utilize read-write locks to allow multiple concurrent reads operations (Member) while ensuring exclusive access for writes operations (Insert and Delete).

The program runs 100 initial runs at the beginning and calculates the mean and standard deviation. Then it calculates the required number of runs to get results within an accuracy of $\pm 5\%$ and 95% confidence level using the following equation,

$$n = \left(\frac{z \times \sigma}{E} \right)^2$$

Where:

- z is the z-score corresponding to the desired confidence level (1.96 for 95% confidence).
- σ is the population standard deviation, which is estimated by the sample standard deviation.
- E is the margin of error.

Then if the required number of runs is greater than the initial number of runs (100), then it will rerun the required number of times.

Performance Evaluation Outcomes:

- The average execution time and standard deviation for each configuration (number of threads) will be saved in the results folder after each run.
- Diagrams will be saved in the diagrams folder after each run.

Computer used:

CPU:

- Intel(R) Core(TM) i7-119G7 @2.90GHz 1.80 GHz
- 4 physical cores, 8 logical processors

Memory:

- size 8GB
- type SODIMM

Operating system:

- Windows 11 64-bit

Tools:

- Compiler: gcc
- Libraries: pthread, sys/time
- Visual studio code

Step 3

Case 1: (mMember = 0.99 | mInsert = 0.005 | mDelete = 0.005)

Implementation	No of threads							
	1		2		4		8	
	Average	Std	Average	Std	Average	Std	Average	Std
Serial	9203.220	294.972						
One mutex for entire list	10878.243	6856.175	16598.130	4345.589	15090.909	4897.508	16014.285	4413.820
Read-Write lock	11110.1836	6754.749	7295.441	7651.378	6039.328	723.983	5922.833	1001.251

Case 2: (mMember = 0.9 | mInsert = 0.05 | mDelete = 0.05)

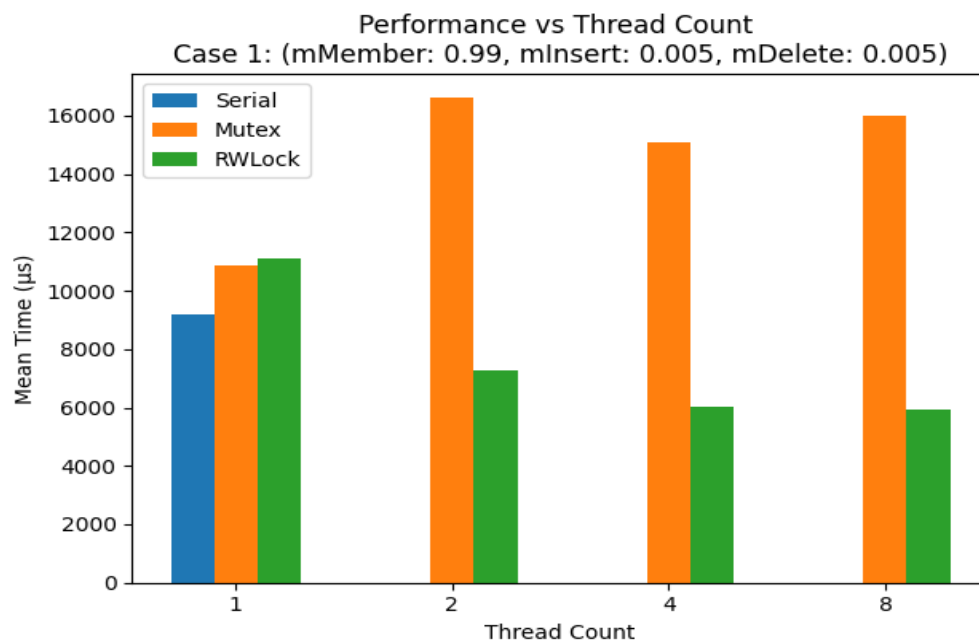
Implementation	No of threads							
	1		2		4		8	
	Average	Std	Average	Std	Average	Std	Average	Std
Serial	13113.390	1034.505						
One mutex for entire list	14199.438	5727.334	18874.172	5855.259	16467.105	5328.814	16550.000	3297.917
Read-Write lock	14875.776	5091.362	16455.357	4608.821	16217.687	5344.972	15410.000	3519.340

Case 3: (mMember = 0.5 | mInsert = 0.25 | mDelete = 0.25)

Implementation	No of threads							
	1		2		4		8	
	Average	Std	Average	Std	Average	Std	Average	Std
Serial	17533.600	812.969						
One mutex for entire list	29236.842	7627.735	31140.000	945.917	20216.748	6924.795	16511.450	4890.125
Read-Write lock	32890.000	6417.927	42900.000	7475.887	33570.000	4841.435	25755.102	7837.905

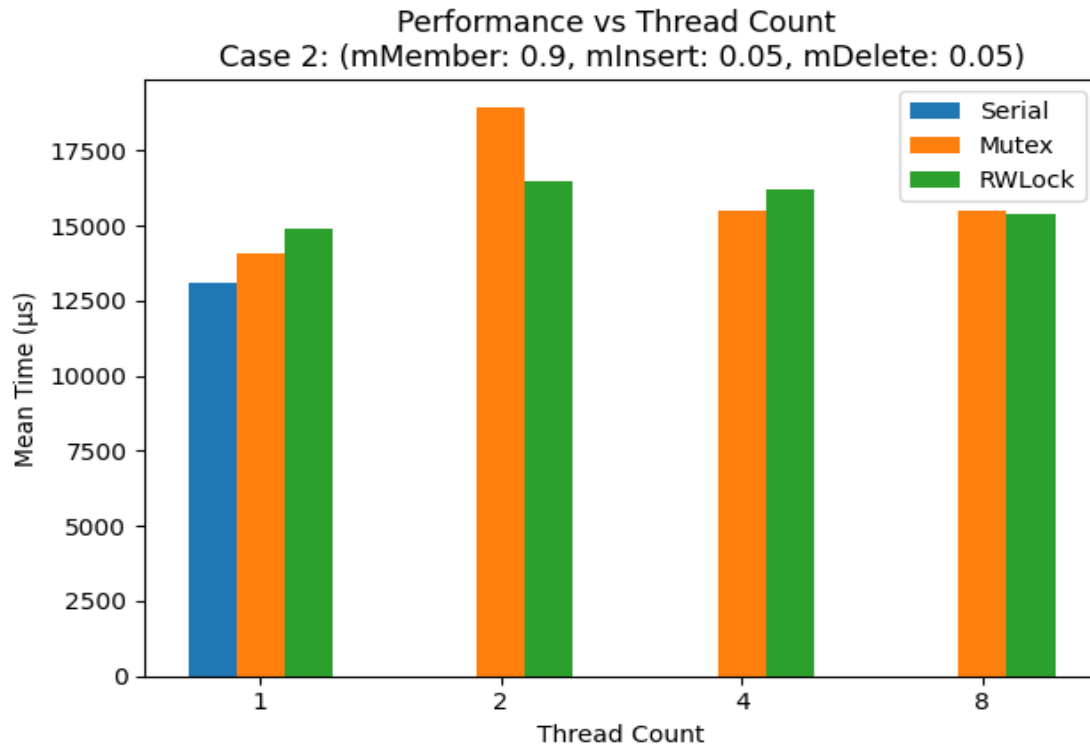
Step 4

Case 1: (mMember = 0.99 | mInsert = 0.005 | mDelete = 0.005)



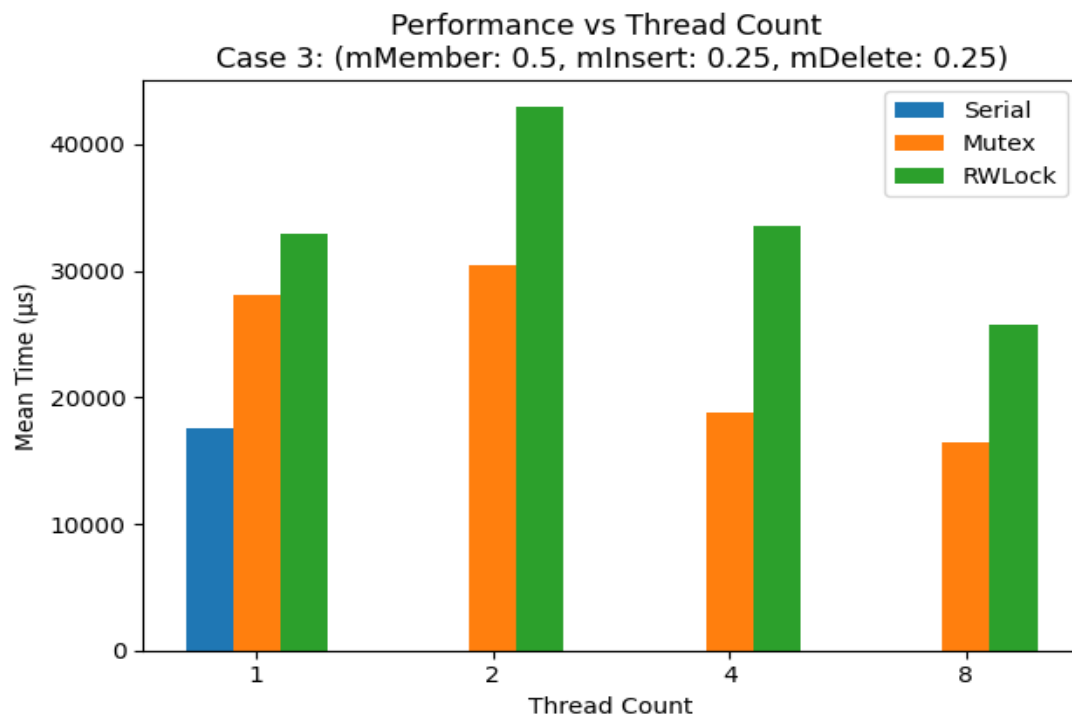
- Diagram 1 -

Case 2: (mMember = 0.9 | mInsert = 0.05 | mDelete = 0.05)



- Diagram 2 -

Case 3: (mMember = 0.5 | mInsert = 0.25 | mDelete = 0.25)



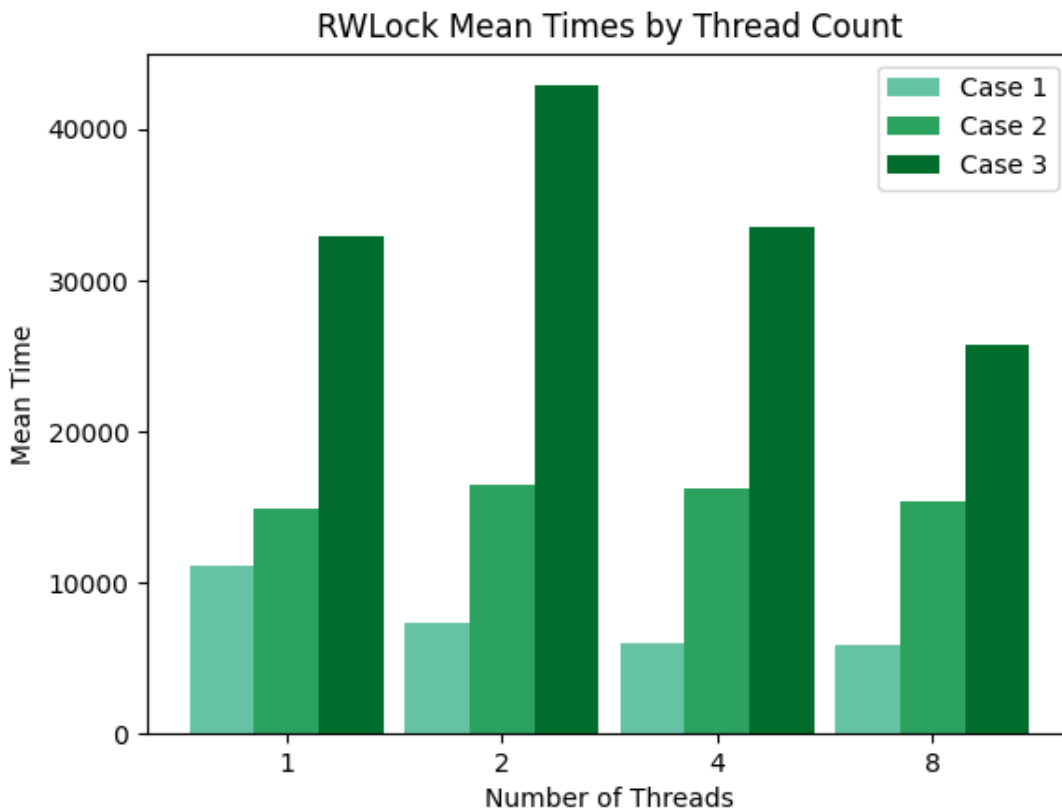
- Diagram 3 -

Step 5

Serial Implementation: This implementation typically performs better than the mutex implementation, due to the absence of locking overhead. RWLock implementation performs better than serial implementation in Case 1, which has more Member operations that can be run in parallel (diagram 1).

Mutex Implementation: In the mutex implementation, the entire linked list is locked for each operation, causing all operations to execute serially, even when using multiple threads. This approach leads to increased execution times due to the high level of contention and the overhead of repeatedly acquiring and releasing the mutex. In situations with higher concurrency (more threads), the performance often deteriorates, indicating the overhead of frequent locking and unlocking operations.

Read-Write Lock Implementation: This implementation generally allows for better read throughput in scenarios with a high read proportion (mMember), as multiple threads can read simultaneously. However, when the number of write operations increases, the read-write lock's efficiency decreases.



- Diagram 4 -

The graph displays the performance of the RWLock implementation across three scenarios with varying thread counts.

- When the number of threads increases, the average running time of the RWLock implementation decreases.
- While RWLock performs optimally in scenarios dominated by read operations (Case 1), its efficiency decreases as the frequency of write operations increases in Cases 2 and 3. This suggests that RWLock is most beneficial in read-heavy environments.